

# **Lab 6 - Neural Networks for Function Approximation**

**Name:** Monisha Sharma

**SRN:** PES2UG23CS906

**Class:** 5F

**Date:** 17/09/2025

## **1. Introduction**

The purpose of this lab was to implement a neural network from scratch to approximate a polynomial curve generated from synthetic data.

Tasks performed:

- Implement key neural network components:
  - Xavier initialization
  - ReLU activation and derivative
  - Mean Squared Error (MSE) loss
  - Forward propagation
  - Backpropagation with gradient descent
- Train the network with early stopping.
- Evaluate performance with training/test loss curves and  $R^2$  score.
- Run hyperparameter experiments and compare results.

## **2. Dataset Description**

**Polynomial Type:** Cubic -  $y = 2.23x^3 - 0.12x^2 + 5.32x + 10.75$

**Samples:** 100,000 total train-80k test-20k

**Noise Level:** Gaussian noise, standard deviation = 10

**Features:** Single input x, single output y

**Scaling:** Both x and y standardized using StandardScaler.

## **3. Methodology**

### **1. Network Architecture:**

- Input (1)
- Hidden Layer 1 (72 neurons, ReLU)
- Hidden Layer 2 (32 neurons, ReLU)
- Output (1, linear).

## 2. Weight Initialization:

Xavier initialization:

$\sim \text{normal dist}(0, \sqrt{2}(\text{fan\_in} + \text{fan\_out}))$

## 3. Forward Pass:

Sequentially compute linear transformations + ReLU, ending with linear output.

## 4. Loss Function:

Mean Squared Error (MSE).

## 5. Backpropagation:

- Compute gradients using chain rule.
- Update weights with gradient descent

# 4. Results

Experiment	Learning Rate	Batch Size	Epochs	Activation	Train Loss	Test Loss	R <sup>2</sup> Score	Observations
Baseline	0.001	64	500	ReLU	0.6255	0.6291	0.3743	Stable, no overfitting, but accuracy is low
Exp 1	0.01	64	300	ReLU	0.2076	0.2091	0.7921	Higher LR gave much faster convergence, best fit
Exp 2	0.001	128	500	ReLU	0.6255	0.6291	0.3743	Larger batch slowed learning, no gain
Exp 3	0.0005	64	500	ReLU	0.8018	0.8059	0.1984	Smaller LR led to underfitting, poor results
Exp 4	0.001	64	800	ReLU	0.5087	0.5120	0.4908	More epochs improved fit compared to baseline

# 5. Conclusion

The baseline model (learning rate = 0.001, batch size = 64, 500 epochs) achieved a moderate performance with a **Test Loss of 0.6291 and an R<sup>2</sup> Score of 0.3743**, showing that the network was able to learn some structure but not fit the data with high accuracy.

Through hyperparameter exploration (Part B):

- A higher learning rate (0.01) significantly improved learning speed and accuracy, giving the best performance (Test Loss = 0.2091, R<sup>2</sup> = 0.7921).
- A smaller learning rate (0.0005) led to underfitting, with poor accuracy (R<sup>2</sup> = 0.1984).

- Increasing the batch size to 128 slowed convergence and did not improve results.
- Increasing the number of epochs to 800 moderately improved generalization ( $R^2 = 0.4908$ ) compared to baseline.

Overall, the best model was achieved with a relatively higher learning rate and smaller batch size, which allowed the network to fit the polynomial more effectively.